

# KUBERNETES WORKSHOP

- Introduction to Kubernetes
- Pods
- Deployments
- Statefulset
- Networking and Services
- RBAC
- Config map & secrets
- Setup -> Local + Kubeadm

# INTRO TO K8S →

- What is Kubernetes?
- Why Kubernetes?
- Kubernetes Architecture
- YAML file
- Kubernetes components
- Imperative vs declarative
- Namespaces
- Labels and Selectors

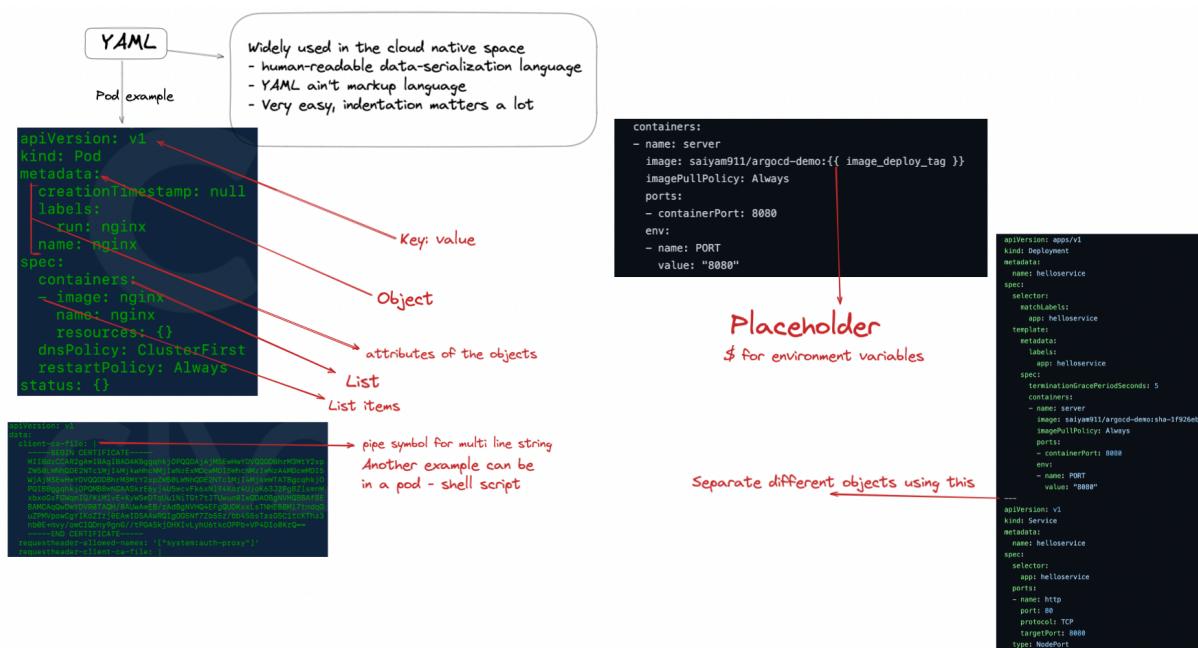
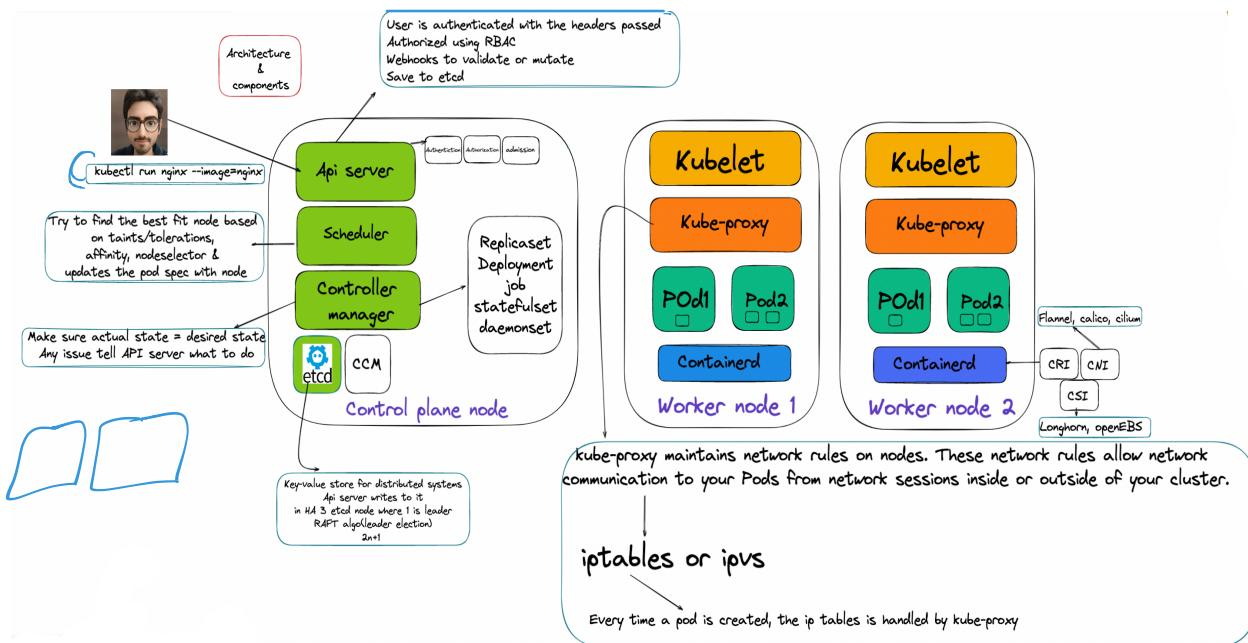
## What is Kubernetes ?

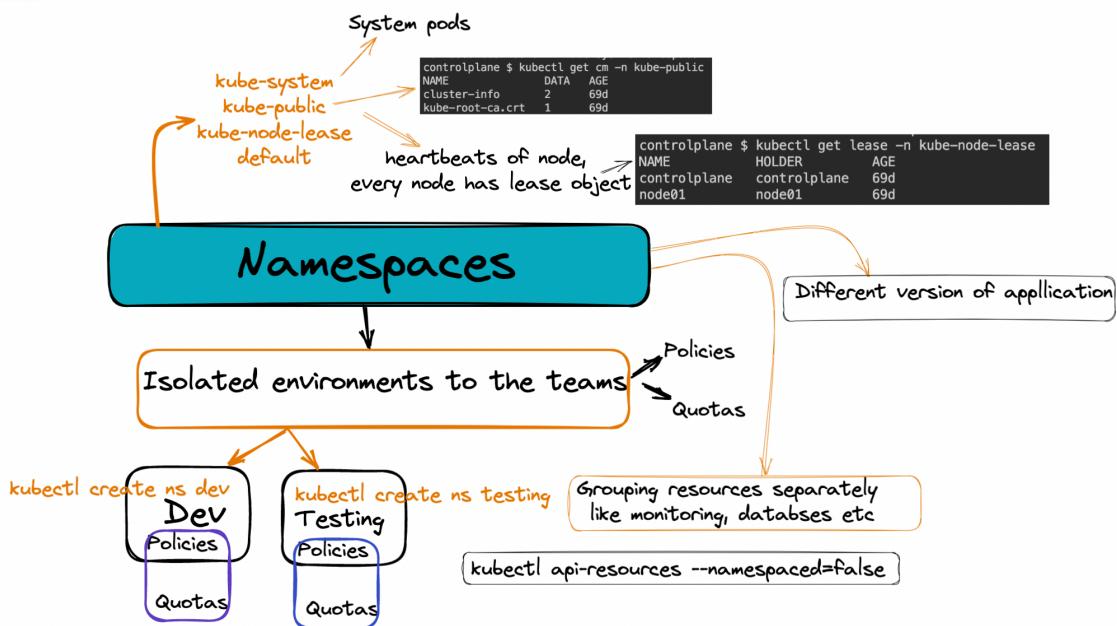
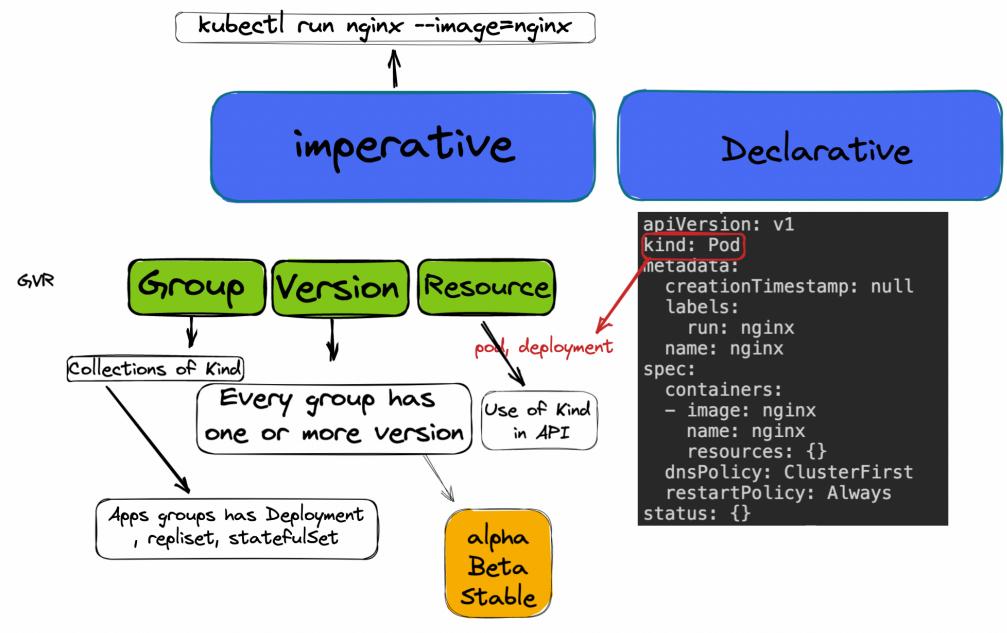
Kubernetes, also known as K8s, is an open-source system for automating deployment, scaling, and management of containerized applications.

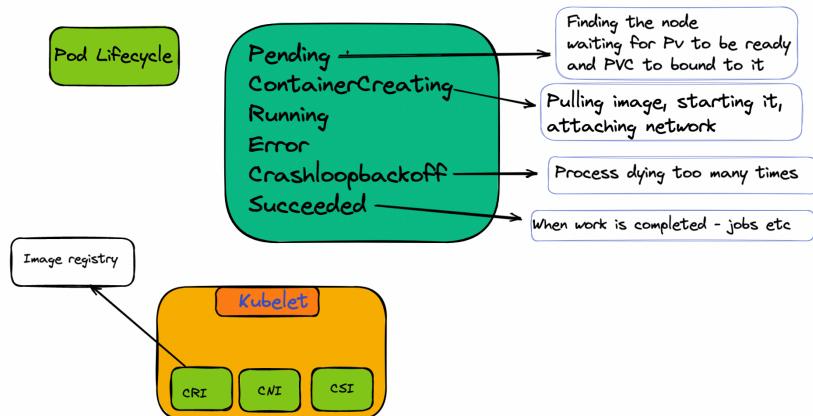
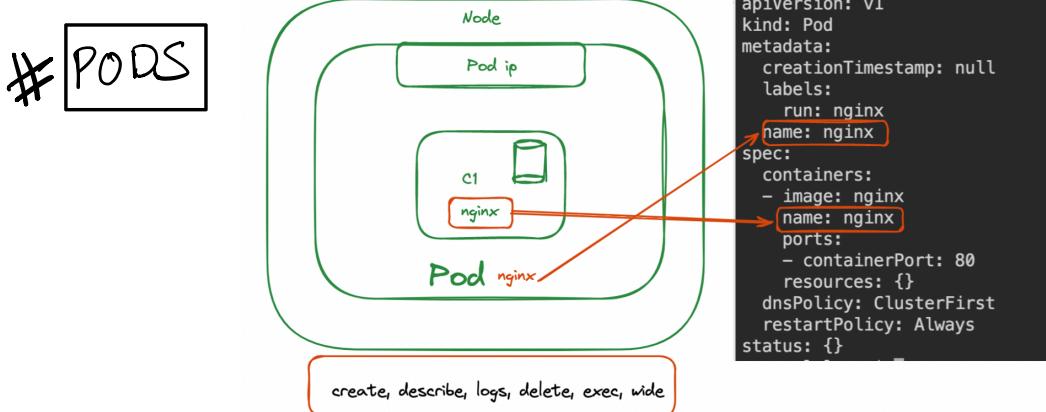
- CNCF Graduated
- Born out of Borg and Omega
- Launched 2014
- Designed for Scale
- Run anywhere

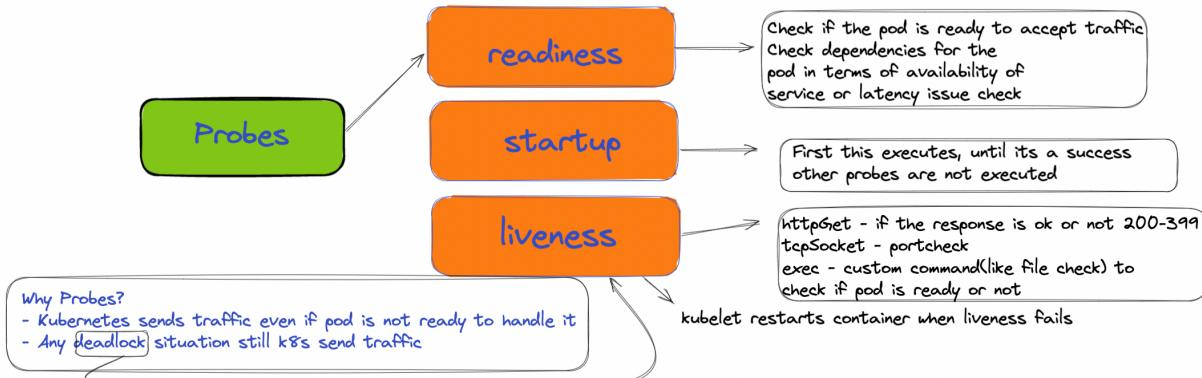
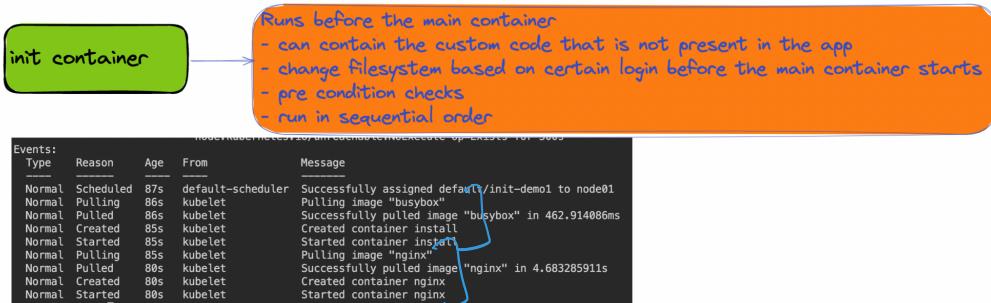
## Why Kubernetes ?

- Few containers are fine but what about scale ?
- How do monitor them ?
- What about running pods on multiple nodes?
- What about flexibility ?
- Autoscaling?
- Scheduling
- Self healing capabilities.



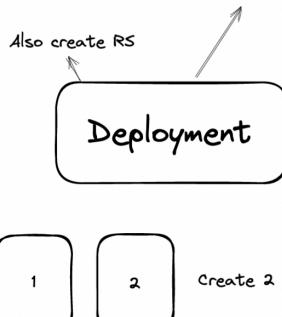






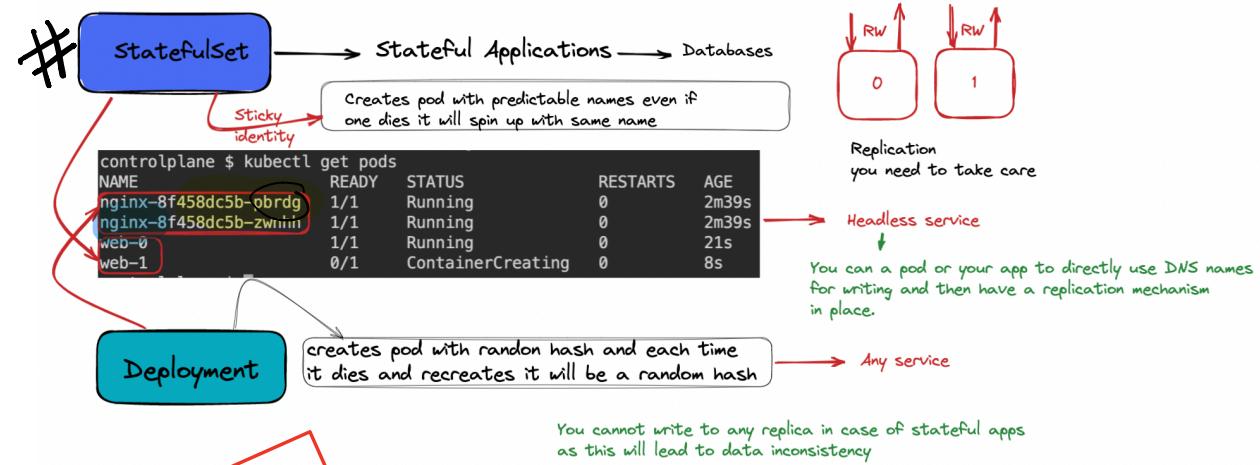
#Deployment

controlplane \$ kubectl get pods  
 NAME READY STATUS RESTARTS AGE  
 nginx-8f458dc5b-8fcf 1/1 Running 0 8s  
 nginx-8f458dc5b-dzphw 1/1 Running 0 8s  
 StrategyType: RollingUpdate  
 MinReadySeconds: 0  
 RollingUpdateStrategy: 25% max unavailable, 25% max surge



```
apiVersion: apps/v1
kind: Deployment
metadata:
  creationTimestamp: null
  labels:
    app: nginx
    name: nginx
spec:
  replicas: 2
  selector:
    matchLabels:
      app: nginx
  strategy: {}
  template:
    metadata:
      creationTimestamp: null
      labels:
        app: nginx
    spec:
      containers:
        - image: nginx
          name: nginx
          resources: {}
status: {}
```

- When you update the deploy image  
 - New pod gets created, old pod serves traffic until it finishes its terminationGracePeriod (30 by default)

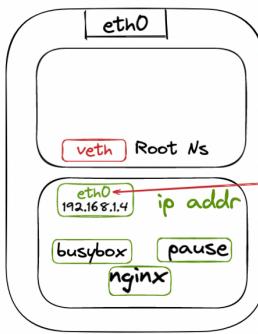


# NETWORKING

What happens when you run the pod?

```
controlplane $ kubectl get nodes
NAME STATUS ROLES AGE VERSION
controlplane Ready control-plane 6hd v1.24.0
node01 Ready <none> 6hd v1.24.0
controlplane $ 
controlplane $ cat pod.yaml
apiVersion: v1
kind: Pod
metadata:
  name: shared-namespace
spec:
  containers:
    - name: p1
      image: busybox
      command: ['"/bin/sh', '-c', 'sleep 10000']
    - name: p2
      image: nginx
controlplane $ kubectl apply -f pod.yaml
pod/shared-namespace created
```

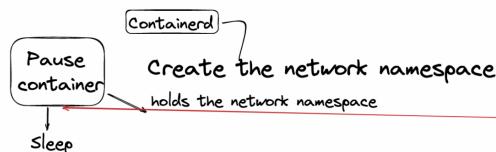
After pod creation CNI attach IP address and attach it to network



```
controlplane $ kubectl exec -it shared-namespace -- sh
Defaulted container "p1" out of: p1, p2
# ip link
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue qlen 1000
  Link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
  inet 127.0.0.1/8 brd 127.255.255.255 scope host lo
    valid_lft forever preferred_lft forever
  inet6 ::1/128 brd 00:00:00:00:00:00 scope host
    valid_lft forever preferred_lft forever
3: eth0@if19: <NOFORWARD,BROADCAST,MULTICAST,UP,LOWER_UP,M-DOWN> mtu 1460 qdisc noqueue
  Link/ether be:26:a7:95:24:17 brd ff:ff:ff:ff:ff:ff
  inet 192.168.1.4/32 brd 192.168.1.4 scope global eth0
    valid_lft forever preferred_lft forever
  inet6 fe80::bc26:a7ff:fe95:2417/64 scope link
    valid_lft forever preferred_lft forever
/ # route
Kernel IP routing table
Destination     Gateway         Genmask        Flags Metric Ref    Use Iface
default         169.254.1.1   0.0.0.0       UG        0      0      0 eth0
169.254.1.1    *             255.255.255.255  UH        0      0      0 eth0
```

**ip netns list**  
List of network namespaces

There are veth pairs created or depending on CNI maybe different

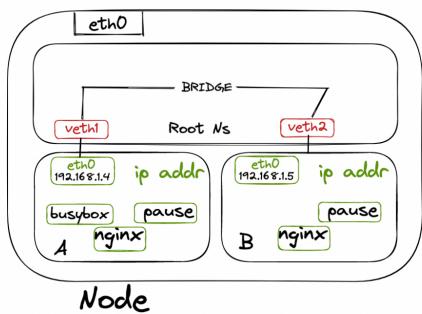


To see that network namespace is held by pause container

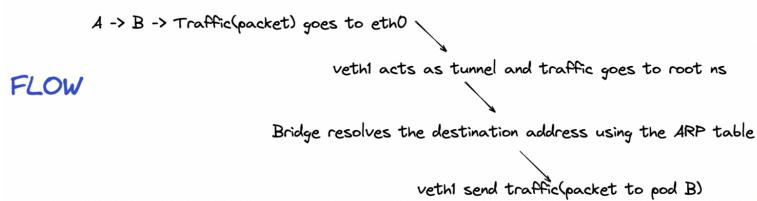
```
node01 $ lsns | grep nginx
4026532987 mnt          2 32059 root           nginx: master process nginx -g daemon off;
4026532988 net          2 32059 root           nginx: master process nginx -g daemon off;
node01 $ lsns -p 32059
NS TYPE      NRPROCS  PID USER COMMAND
4026531835 cgroup     157   1 root /sbin/init
4026531837 user      157   1 root /sbin/init
4026532518 net          3 31073 65535 /pause
4026532584 uts          3 31973 65535 /pause
4026532586 ipc          3 31973 65535 /pause
4026532587 mnt          2 32059 root nginx: master process nginx -g daemon off;
4026532588 pid          2 32059 root nginx: master process nginx -g daemon off;
```

**ls -lt /var/run/netns** → List of namespaces created

```
node01 $ ls -lt /var/run/netns
total 0
-r--r--r-- 1 root root 0 Jul 16 14:25 cni-3cf655f5-9131-e23d-0756-b7ab6556ef8f
-r--r--r-- 1 root root 0 May  8 19:46 cni-06382f28-d5db-63f3-00ee-654e40ee904f
-r--r--r-- 1 root root 0 May  8 19:46 cni-592cc23b-1d2b-6054-bf32-437fd23a04df
node01 $ ip netns exec cni-3cf655f5-9131-e23d-0756-b7ab6556ef8f ip link
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN mode DEFAULT group default qlen 1000
  link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
3: eth0@if9: <NOFORWARD,BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1460 qdisc noqueue state UP mode DEFAULT group default
  link/ether 9e:bd:c7:d4:e5:99 brd ff:ff:ff:ff:ff:ff link-netnsid 0
node01 $ ip link | grep -A1 9
9: calic40f455693@if3: <NOFORWARD,BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1460 qdisc noqueue state UP mode DEFAULT group default
  link/ether ee:ee:ee:ee:ee:ee brd ff:ff:ff:ff:ff:ff link-netnsid 1
node01 $
```

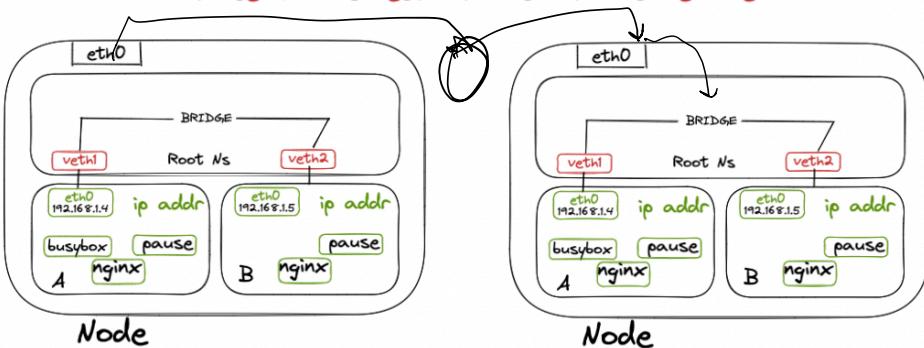


### Inter node pod communication



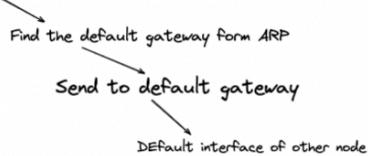
### Node to node communication

Cannot use ARP as destination is not on same node



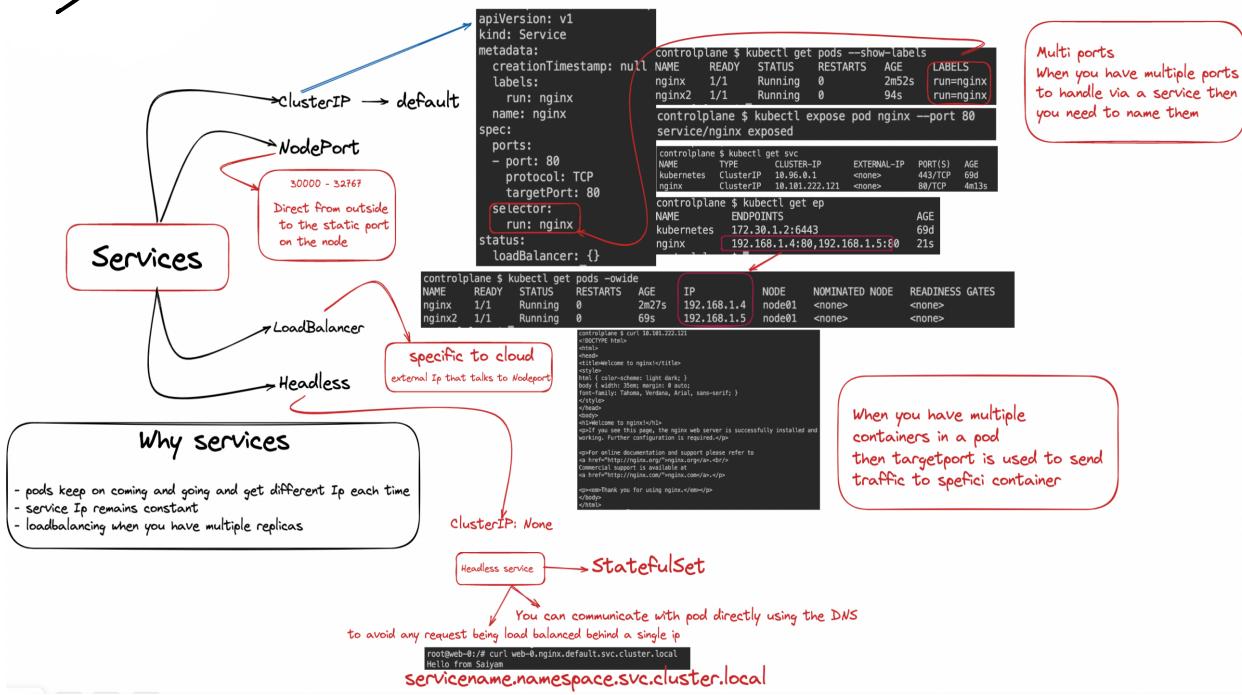
### FLOW

To check if the destination is not on the same network bitwise or Adding operation is performed

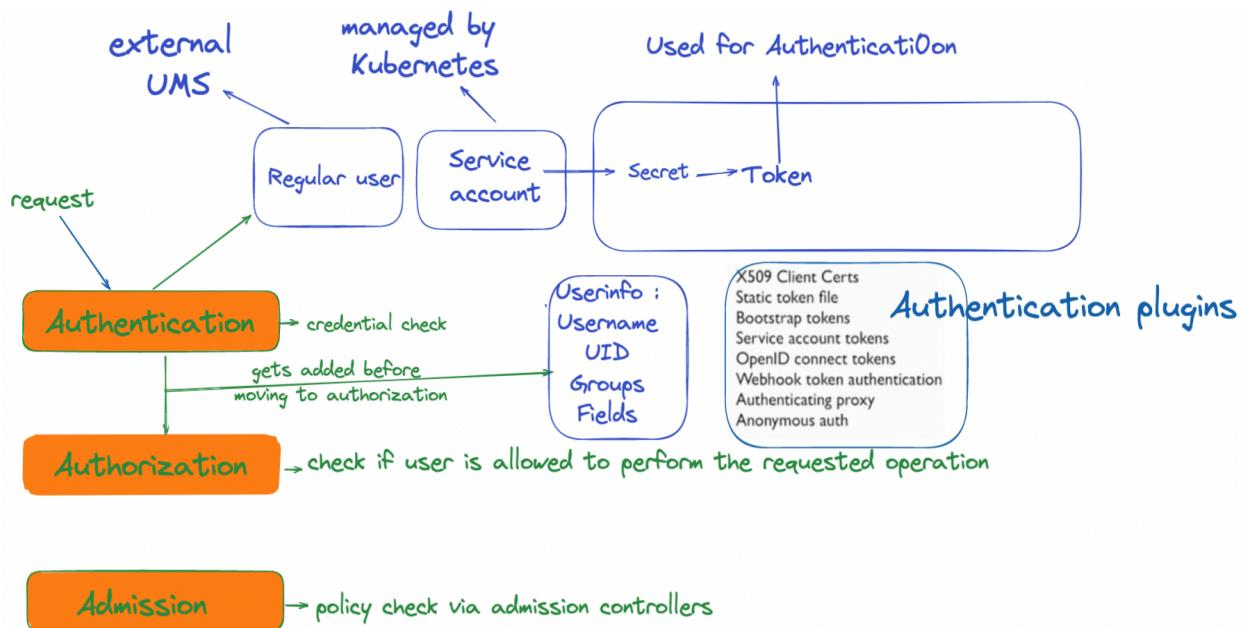


Services — iptables and netfilter

# SERVICES



# # AAA



*Config map  
# secrets*

